Functional Queues

Kristoffer Just Andersen, 20051234

Troels Leth Jensen, 20051234

Morten Krogh-Jespersen, 20022362

AARHUS
UNIVERSITY
DEPARTMENT OF COMPUTER SCIENCE

Contents

Chapter 1

Introduction

This is the third and final report in the Advanced Data Structures course at Aarhus University and we will in this project investigate the performance of queues in the functional language by the name of Haskell.

We first present some terminology and then a total of five different queues where two only differs by one being lazy. Hereafter, we describe how we made the implementation and tested for correctness and why we can say that our implementations are correct.

We then describe how we measure running times in a lazily evaluated language and explain the different tests we have chosen for testing the running times of each queue.

Finally, we will give an answer to the problem from the theoretic project of how to implement a data structure having push and pop in $O(1)$ while also having index lookup in $O(\log n)$ and conclude on the entire project.

1.1 Terminology

We will use some specific phrases to describe properties that we, for good measure, include below:

Purely functional   Purely functional is a term that is used to describe data structures where no operation perform destructive updates. Such a data structure is said to be immutable or persistent. Notice that, in this course, we further divide persistence into partial persistence and full persistence, where this is considered to be full persistence.

Strict evaluation   Strict evaluation (eager evaluation) is where an expression is computed at the point where it is assigned to a variable or is passed as an argument.

Lazy evaluation   Lazy evaluation delays the computation of a variable such that the expression is computed at the point where the variable is used.

Memoization   When an expression is bound to a variable, that variable can be used multiple times. With a strict evaluation, the variable is assigned a specific value. If a lazy evaluation strategy is used then the first time the value of a variable is needed, the computation will occur, but what about subsequent uses? Memoization says that when a value has been calculated once, that value will be stored and subsequent uses will therefore not repeat any calculations. The combination of lazy evaluation with memoization is also called call-by-need.

Chapter 2

# A purely functional list as a queue

Figure 2.1: A purely functional list

Figure 2.1 shows a graphically representation of a functional list where elements can be inserted at the front by means of **cons** and the list can be traversed by following the arrows which corresponds to the **cdr** operation. The list can be constructed in constant time by a special termination construction (inductive base case) being the empty list.

Elements can be inserted and removed from the front in constant time, but all other operations require traversing the entire list. Below are all operations with the corresponding time-complexities shown:

Operation	List
MAKELIST(x)	$O(1)$
PUSH(x,L)	$O(1)$
POP(L)	$O(1)$
INJECT(x,L)	$O(n)$
EJECT(L)	$O(n)$

Because a queue pops elements from the front and injects elements in the end, this is not the most efficient queue implementation.

Chapter 3

Queue by a pair of lists

A simple observation to make is that if we reverse the list of section 2, $\text{INJECT}(x, L)$ and $\text{EJECT}(L)$ will be at the front of the list and will therefore take $\text{O}(1)$ time. However, $\text{PUSH}(x, L)$ and $\text{POP}(L)$ will take $\text{O}(n)$ time. A trick is therefore to maintain a pair of lists such that $\text{PUSH}(x, L)$, $\text{POP}(L)$, $\text{INJECT}(x, L)$ and $\text{EJECT}(L)$ for the most part perform in constant time.


Figure 3.1: Maintaining two lists - reverse B and place as A on next pop from A

When we enqueue nodes they are pushed to the front of list B, and when elements are removed they are taken from the front of list A. Both of these operations will normally take constant time. The only problem is what to do if A is empty.

Figure 8.1 shows a configuration of the queue where a pop will result in A becoming the empty list. Subsequent pops should remove from the end of B but that would take $\text{O}(n)$ for every operation. It therefore makes more sense to reverse the list once and place it instead of A.

To analyze the running time we will use amortization and specifically use the Potential method [1, p. 459]. We will use the length of the tail list as the potential function. Every insert into tail is therefore an operation that takes a single step and increases the potential by one, so the amortized cost is two.

Reversing the list can be done in $\text{O}(|tail|)$. The length of the tail is the number of enqueues since the last reversal or construction of the list. Therefore, we must have $|tail|$ unreleased potential, giving us that the tail list can be reversed in $\text{O}(1)$ whenever needed. The queue will have the following time complexities:

Operation	List	Two lists amortized / worst case
$\text{MAKELIST}(x)$	$\text{O}(1)$	$\text{O}(1)$ / $\text{O}(1)$
$\text{POP}(L)$	$\text{O}(1)$	$\text{O}(1)$ / $\text{O}(n)$
$\text{INJECT}(x,L)$	$\text{O}(n)$	$\text{O}(1)$ / $\text{O}(1)$

3.1 Converting to strict in Haskell

Laziness is a feature in Haskell. A queue might perform differently depending on the choice of evaluation. At the very least, the expensive operations will be performed at different runtime points in the program. By placing strictness annotations the right places, (putting ! in front of the two constituent lists) we tell Haskell to enforce a strict evaluation of the data structure. That is, a value representing this type of queue will never contain any *thunks* or unevaluated values. Because of the way rotations happen only when needed, we might not see any differences between lazy or eager evaluation for this queue. Nevertheless, we will perform tests on both the lazy and the strict version of this queue. We will delay (are we then lazy?) the discussion about amortized running times for a lazy version of this queue to the next chapter.

Chapter 4

Pair of lists with lazy evaluation

We were asked to implement a queue with lazy evaluation having amortized constant time for $\textsc{inject}(x, L)$ and $\textsc{pop}(L)$. We already have one such implementation but it might be fun to have another where rotation occurs at points in execution where the list reversion can be delayed.

Like other functional queues we have two lists; one representing the front and another representing the tail. At specific points we will rotate the elements from the tail onto the head by reversing the tail list and appending it. In Haskell the append and the reverse function is lazy, giving us the desired evaluation strategy.

The trick with two lists has already been used in the previous chapter but we will use another trigger for the lazy implementation. We rotate whenever the length of the tail list is one greater than the length of the front list.

4.1   Analyzing lazy evaluations

When we calculated the potential in the previous section we made the assumption that a queue was used ephemerally. We did not discuss amortized running times when using the *same* queue multiple times. Consider a queue that has a lot of potential stored and the triggering of a reversal of the tail list is imminent. The queue is persistent, meaning, that performing $\textsc{pop}(L)$ on the same queue is non-destructive. Therefore, the expensive operation can be computed multiple times if we pass the original queue as argument to different operations. The released potential can only be used once, thereby ruining our amortized analysis.

There is no way to avoid this computation in language having call-by-value or call-by-name without using side effects, but for lazy evaluations with memoization, which we call call-by-need, the expensive operation is only carried out once. However, memoization is a side effect and thus, some of our arguing might be considered a test of faith.

It is inadequate to do amortized analysis for lazy amortized data structures by arguing as if they were strict [3]. We divide the cost of an operation two cat-

egories; unshared cost and shared cost. Unshared cost is the cost the operation takes, with the assumption that all lazy constructions of the data structure, that will be used, have been computed. The shared cost of an operation is the time it would take to execute every suspension created but not evaluated by the operation. The sum of both is called the complete cost.

To adjust he Potential function we have to replace the savings function Φ (we store potential) with a function Ψ representing the upper bound on accumulated debt. The amortized cost of an operation will therefore be the complete cost of an operation - the change in potential. With this change we can calculate the complete cost of an operation as if it is strict.

4.2 Completing the analysis of the lazy queue

We will analyze the lazy queue with our modified potential function Ψ. As always, it should be the case that we add potential on inserts and release potential when deleting. We therefore define:

$$\Psi(Q) = |notreversed(F)| + |R|$$

Let us analyze the amortized cost of an INJECT(x, L) operation and let us consider the case where no rotation occurs. The operation takes one step and also has to pay to the potential function, which is a total cost of 2. The enqueue that creates a rotation also takes one step when adding and also have to pay the cost of increasing the potential. When the rotation occurs the potential described by $|R|$ will be released by the same amount that $|notreversed(F)|$ increases resulting in no additional change in potential. Finally, we have to remember to pay for initiating the rotation which gives a total of 3.

Analyzing POP(L) is straight forward, however there are four cases. The case where POP(L) does not force an unevaluated expression only has the cost of removing the first element. If POP(L) forces a rotation, the cost is 2 per the argument above. If POP(L) evaluates an unevaluated expression, the operation will take O(n), but the potential will decrease by the same amount, giving us a total cost of 2. Finally, if POP(L) rotates and forces an unevaluated expression to be computed, the final cost will be 3.

The above analysis assumes that, when a demand for the first element of a lazy list not yet reversed, the list is reversed and all subsequent calls to the list will work in constant time. The latter part assumes therefore that lazy lists are memoized. We can now present the time complexities for the lazy algorithm:

Operation	List	Two lists amortized / worst case	Lazy amortized / worst case
MAKELIST(x)	O(1)	O(1) / O(1)	O(1) / O(1)
POP(L)	O(1)	O(1) / O(n)	O(1) / O(n)
INJECT(x,L)	O(n)	O(1) / O(1)	O(1) / O(1)

Chapter 5

Hood-Melville - $\mathrm{O}\left(1\right)$ worst case

All of the queues described above have worst case running times $\mathrm{O}\left(n\right)$, which is unacceptable in some situations. Thus there exists a need for real times queues where all operations have worst case running times $\mathrm{O}\left(1\right)$. We have chosen to implement the Hood-Melville real time queue [2], in the style of [3].

There are a few key insights to this algorithm. First, reversing a list incrementally can be done by having two lists and transferring elements from one to the other. Second, one can incrementally append two lists by applying the trick. If one would like to append two lists xs and ys, one can reverse xs to xs' and then reverse xs' onto ys. Additionally, one can reverse xs on to the reverse of ys by reversing xs and ys in parallel and then continue as before.

Initially, the queue will have two lists forming the head H and tail T and two integers describing the length of each. When we decide to rotate, we use the observation above and do the following:

1. Reverse T forming the tail of the new resulting head H'

2. Reverse H into H_R

3. Reverse H_R onto H'

One should be able to convince him- or herself that queue order is preserved after the third step. However, the queue will not remain constant when performing the rotation (which is a weird phrase in [2] since the incremental rotation only occurs when changes occur), but basically, it means that we have to maintain the state of the current rotation and provide the user the ability to still make alterations to the queue.

Allowing the user to enqueue elements is easily done by just having a new tail list. Dequeuing is bit harder because we are currently reversing the H into H_R. The answer is to have two lists: one being a working copy of the old list and one being the list we are reversing. This again introduces some maintenance because removing an element from the working copy somehow has to be synchronized with the list we are reversing. To correct this, we use a

counter that describes how many valid elements to be copied from H_R. A total of six lists is therefore necessary for this data structure.

The copying (or rotation) should be completed before the first element is needed. We will rotate elements when the tail list becomes one longer than the head list. Therefore, rotating elements will take $m + 1$ operations to complete Step 1 and m operations to complete Step 2 Since these can be done in parallel the total time is $m + 1$. Completing Step 3 will take additional m operations giving a total of $2m + 1$. Therefore $2m + 1$ incremental operations must be performed in at most m queue operations. By performing the two first rotation steps starting the rotation process and hereafter perform two incremental steps for each queue operation we should be on the safe side. In total this gives $2(m + 1) = 2m + 2$ steps, which is greater than $2m + 1$. We therefore perform constant work for every queue operation.

Operation	List	Two lists amortized/worst	Lazy amortized/worst	Hood Melville
MAKELIST(x)	$O\left(1\right)$	$O\left(1\right)$ / $O\left(1\right)$	$O\left(1\right)$ / $O\left(1\right)$	$O\left(1\right)$
POP(L)	$O\left(1\right)$	$O\left(1\right)$ / $O\left(n\right)$	$O\left(1\right)$ / $O\left(n\right)$	$O\left(1\right)$
INJECT(x,L)	$O\left(1\right)$	$O\left(1\right)$ / $O\left(1\right)$	$O\left(1\right)$ / $O\left(1\right)$	$O\left(1\right)$

Chapter 6

Testing and correctness

Testing a data structure against a specification is traditionally done using some permutation of black-box and unit testing, as done in the previous reports. Implementations in C were tested against hit-and-miss unit tests at the whim of the programmers, to the extent that their imagination and endurance permitted.

We can do one better with the levels of abstraction permitted by a functional language – in particular the type system of Haskell.

We test the implementations by using a concept from operational semantics known as *observational equivalence*, as detailed by John Hughes & Koen Claessen [**?**]. It uses the idea of an *evaluation context*, a notion of "the surrounding program", and of *observable results*, the effects visible to the surrounding program caused by running a particular piece of code. Two pieces of code are then observationally equivalent if they produce the same observable result in all evaluation contexts.

We here, like the paper, restrict ourselves to "queue programs" – the subset of Haskell dealing only with our queue interface. While we, for correctness, should consider all of Haskell as evaluation contexts, we are interested in the behaviour of our implementations, and – as opposed to C – there is no way to manipulate the state of the queue without going through the interface. Without pointer arithmetic, side effects and associated suspects, we can be certain of the absence inadvertent mutation of the international structure of the queue.

A queue is an inherently "stateful" object, and so is typically used in an imperative fashion. As such we consider imperative-style monadic programs of Haskell as the context, which only manipulates the queue through the interface, as argued above. We define a data type for representing such contexts:

```
data Action = Inject Int
            | Pop
            | Peak
            | Return (Maybe Int)
```

representing an inject, a pop, a peak, and the binding of the result of a peak, respectively. This last action is needed to argue that instead of obtaining an element from a queue, if we know that element, we might as well explicitly write that element in the program ourselves.

We then define a function "executing" an evaluation context on a given queue, collecting observables as it goes. It reads equationally as one would expect.

```
perform :: Queue q => q Int -> [Action] -> [Maybe Int]
perform q []      = []
perform q (a:as) =
  case a of
    Inject  n -> perform (inject n q) as
    Pop       -> perform (pop q) as
    Peak      -> peak q : perform q as
    Return m -> m : perform q as
```

Finally, we define a testable property on queues, that asserts that in all evaluation contexts (prefixes and suffixes of actions on queues) we observe the same results from evaluating the queue in that context.

```
equiv :: Queue q => q Int -> [Action] -> [Action] -> Property
equiv q c c' =
  forAll (actions 0) $ \pref ->
  forAll (actions (delta (pref ++ c))) $ \suff ->
  let
    observe x =
      perform q (pref ++ x ++ suff)
  in
   observe c == observe c'
```

Finally, we wire the property into the QuickCheck library [**?**] which generates random evaluation contexts, and verifies 5 assertions on queues, which according to Hughes & Claeseen [**?**] defines a suitable specification of queues:

```
prop_PeakInject q m n =
  equiv q [ Inject  m, Inject  n, Peak] [ Inject  m, Peak, Inject  n]
prop_InjectPop q m n =
  equiv q [ Inject  m, Inject  n, Pop] [ Inject  m, Pop, Inject  n]
prop_PeakEmpty q =
  equivFromEmpty q [Peak] [Return Nothing]

prop_PeakInjectEmpty q m =
  equivFromEmpty q [Inject m, Peak] [ Inject  m, Return (Just m)]
prop_InjectPopEmpty q m =
  equivFromEmpty q [Inject m, Pop] []
```

The essence of the argument for the specification is to "bubble" peeks and pops towards the beginning of the program with the first 3 properties, and then eliminate them with the 2 following properties, thus "normalizing" programs on queues.

Finally we run the test for each of our 4 implementations, creating a hundred programs of an average length of 8 for each property for a total of $4\cdot5\cdot100 = 2000$ unit tests. These are just default settings and as good as any, but more tests naturally strengthen the claim that the implementations meet the specifications.

```
runTests :: IO ()
runTests = do test "SimpleQueue" (empty :: BasicQueue Int)
```

```
        test  "SmartListQueue" (empty :: PairQueue Int)
        test  "OkasakiQueue" (empty :: OkasakiQueue Int)
        test  "RealTimeQueue" (empty :: RealTimeQueue Int)
where
  test s q = do putStrLn $ "### Testing " ++ s ++ " ###"
               suite q
```

It is unfortunate that the tests do not conclusively prove the correctness of our implementations. If only there was a way...

It turns out, there is! As all our data structures are purely functional and do not use general recursion in the operations, we can implement the data structures in the Proof Assistant Coq. Which we did. Coq allows us to state properties on functional programs, *conclusively prove them*, and then extract the functional programs to e.g. Haskell. As an added benefit, just by virtue of using Coq as an implementation we certify our functions total and terminating, as Coq does not accept partially defined functions.

We start by giving an abstract module declaration of a queue implementaion.

```
Module Type QUEUE.
  Parameter Q          :                 Type −> Type.
  Parameter inv        :  forall  A, Q A −> Prop.
  Parameter empty   :  forall  A, Q A.
  Parameter inject   :  forall  A, A    −> Q A −> Q A.
  Parameter pop      :  forall  A, Q A −> Q A.
  Parameter peak     :  forall  A, Q A −> option A.
  Parameter to_list  :  forall  A, Q A −> list A.
```

The inv field is a predicate on a queue representing whatever invariants might be applicable for a given implementation, and the to_list function is used to convert a queue to an ordered list of elements, which we shall use as our model for specifying queues. Using an abstract model (ordered lists) and then showing our expectations in the model are satisfied by the implementation is much in the spirit of observational equivalence: if looks like we expect in all ways we can measure it, it must be the same.

Next, we define a series of proof obligations, to describe what must be proven about a given implementation. Among them are that the invariant is satisfied by an empty queue, and preserved by all operations. Do not be alarmed by the use of "Axiom": it means "to be proven", not "to be assumed".

```
Definition  represents  {A : Type} (q : Q A) (l :  list  A) : Prop :=
   to_list  _ q = l.

Axiom empty_inv :
   forall   A,
     inv  A (empty A).

Axiom inject_inv :
   forall   A q x,
     inv  A q −>
     inv  _ ( inject  _ x q).
```

```
Axiom pop_inv :
  forall A q,
    inv A q −>
    inv _ (pop _ q).
```

The represents field is used to tie a given instance of a queue to its representation in the model. With that we can define the following proof obligations:

```
Axiom empty_is_empty : forall A,
  represents (empty A) nil.
```

The above dictates that an empty queue represents the empty list of elements.

```
Axiom inject_spec :
  forall A l q x,
    represents q l −>
    inv _ q −>
    represents ( inject A x q) (l ++ (x :: nil )).
```

The above states that injecting an element corresponds to postpending an element to the ordered list.

```
Axiom pop_empty :
  forall A q,
    represents q nil −>
    inv _ q −>
    represents (pop A q) nil.
```

The above states that that popping the empty queue returns the empty queue.

```
Axiom pop_xs :
  forall A q x xs,
    represents q (x :: xs) −>
    inv _ q −>
    represents (pop A q) xs.
```

The above states that popping an element from a queue which represents the list starting with x followed by xs yields a queue representing xs.

```
Axiom peak_empty_spec :
  forall A q,
    represents q nil −>
    inv _ q −>
    peak A q = None.
```

The above states that peeking an element from an empty queue yields nothing.

```
Axiom peak_spec :
  forall A q x xs,
    represents q (x :: xs) −>
    inv _ q −>
    peak A q = Some x.
```

```
End QUEUE.
```

Finally, we state that peeking from a queue representing the ordered sequence starting with x yields that x.

The basic, naive functional queue is an implementation of our model, and it follows by construction that the proofs are largely uninteresting. So are the naive implementation based on a pair of queues. We wont give a detailed account of the proofs of each implementation, as the proof scripts are quite monstrous to present in writing. We refer curious readers to Queue.v in the accompanying code.

Finally, when all proof obligations are satisfied we run e.g.

 Extraction "PairQueue.hs" PairQueue

to generate a Haskell file with the definitions given in the module. While slightly unidiomatic in style, the generated code is exactly as we would have handwritten, and so suffer no performance penalties.

The generated files are exactly those used in our benchmarks and the previously outlined random testing, so it is no surprise it satisfied the previous specification!

We might remark that Coq is a strict programming language which we use to reason about programs to be run in a lazy setting. This is inherently not a problem as none of the implementations exploit lazy capabilities of lazy evaluation but use it only to argue performance characteristics. We have just proven they have the desired semantics and so will behave the same in a lazy setting, with the added benefit of "lazy performance".

At the time of writing, we have not succeeded in proving the Hood-Melville real-time queue correct. Instead, we have made headway in establishing a useful invariant by proving that inject preserves the empty invariant, and the empty queue establishes it – all under the assumption that it is the right invariant, of course. Coq still allows extraction to Haskell, however, with loud warnings about running uncertified code. We believe it doable with more time and effort invested.

Chapter 7

Time measurements

We have a total of four conceptually different queues - where two comes in the two flavours, lazy or eager evaluated, for a total of six. We are asked to design benchmarks where a queue is only used once and also where the queue can be used multiple times. In this chapter we describe the different tests that we apply to the queues and we discuss the results of said tests. Per default, Haskell will evalute to first seen constructor, however, when doing tests this might give us some strange timing results. Therefore, we make sure that the queue under test is evaluated to normal form.

7.1   One punch tests

We have designed two tests that enable us to directly measure time of insertions and insertion + deletion. However, a queue is useful for more things than theoretic time analysis so we are also using the queue to do Breadth-First search in graphs to have a sort of real-life example.

7.1.1   Time of insert

This tests is analogue to the insert test we performed in Project 2. We would like to investigate, on average, how long it takes for a number of insertions to be inserted depending on the size of the queue. We fill up the queue to some amount x and then perform a constant number of inserts. The total time is then divided by the constant and we have a measure for how long a single insert takes depending on the size of the queue.

7.1.2   Time of insert + delete

Rotation will not occur on the same points in execution based on test size, thus if we were to test time of delete by just deleting a few items when the queue has x elements, we might force a rotation for one queue but not another. To counter this obstacle we decided to measure the total time of inserting n elements and removing all elements again.

7.1.3 Breadth-First Search

A classic application of queues in functional programming is in a breadth first traversal of a tree. We made a test case that traversed a complete balanced binary tree of increasing depths. The case exercises the queue with an even mix of injects and pops, where every element is enqueued and popped exactly once for a total of 2^d operations for a tree of depth d.

7.2 Rocky tests (multiple punches - movie reference)

In this section we describe tests that perform more than one operation on the same queue. These experiments are necessary because of the persistent nature of queues, which is clearly explained in section 4.1.

7.2.1 Insert onto the same queue

This test is designed to hit worst cases on all our queue implementations, if hitting such a thing is possible. As in the single usage case, we fill the queue up to some length n. Herefter, we delete and insert c times on the queue given as argument, insert an element such that the size is increment by one $(n + 1)$ and the iteration is repeated. We stop when we have performed the $2 \cdot n$ iterations.

7.2.2 Queue-sort

We define a version of Queue-Sort that is not very efficient. Given n queues of size n where elements are ordered by apperance in the queue, merge all elements into a single queue.

The elements of the queue is generated at random with numbers between the last element inserted in a list q and $|q| \times M/n$.

For each iteration, a fold is applied to the list of queues where the accumulating information is the currently smallest element, the tail of the queue the pop was performed on and the index of the q in the list. The result is stored in a local variable. Hereafter, a map is performed on the same list, where all queues are passed along, except the one that had the smallest element, which is replaced by the tail of the previous queue. The mapped result is then given as argument to the next iteration.

Performed on different initial lengths of the queues, we hope to see a number of expensive operations performed multiple times on the same queue.

7.3 Results

In this section we present the results of our experiments. All running times has been measured and compared to a meanThe running time for the List grows much faster than the others, making it harder for us to distinguish between the rest of the queues. Therefore, the List running times are from all of our graphs, but the original graphs can be found in the appendix.

7.3.1 Time of insert

This test was supposed to measure the time a single insertion takes depending on the fill-rate and the time usage is plotted below:


Figure 7.1: Time per insert operation when queue is of size n

The results are completely as expected. There is only a slight overhead when cons'ing elements on a list and all takes a constant time which matches the theoretic bounds.

7.3.2 Time of insert + delete

This test measures how quickly a queue can be filled and emptied again:

Figure 7.2: Average time per operation for filling queue with n elements and emptying it again

Again we can conclude that the theoretic bounds are met for deletion. Additionally, we see, that the Hood-Melville takes a bit longer per operation than the others. It was expected that maintaining the rotation would have a constant factor of penalty.

7.3.3 Results of Breadth-First Search

Below is the graph-plot of the Breadth-First Search results:

Figure 7.3: Total time of Breadt-First search for a complete binary tree divided by the number of nodes in the tree

Interestingly, for larger levels, there is not much difference in the time per operation. This was not completely expected but since every node must be popped from the queue, every queue forces a complete evaluation.

For the Pairs queues there is exactly one rotation per level, which is easy to convince onself about, because when a level is completed, the tail list will contain all childs for the next level and the head list will be empty.

For the Okasaki queues there will be two rotations for each level. We will give a simple argument. In the very best case, beginning on a new level, all elements n will be in the head list H. This is the farthest from a rotation we can be. Since every dequeue will add two new elements to the tail a rotation will occur in the $n/2$ step when adding the left child to the queue. After rotation and the completed $n/2$ steps we have $|H| = n - n/2 + n - 1 = 3/2n - 1$ and $|T| = 1$, with the remaining $n/2$ nodes of the level yet to be dequeued. After $n/2$ dequeues $|H| = n - 1$ if no rotation happens but the tail would be $|T| = n + 1$ and another rotation would therefore have happened. There is no way of avoiding at least two rotations per level.

But the penalty of rotating $n/2$ elements two times instead of n elements one time is not significant, therefore we see similar results.

7.3.4 Insert onto the same queue

This is the first test that tries to force multiple expensive operations for each queue and the time-usage is shown below:

Figure 7.4: Time of performing multiple inserts and deletes on the same queue

We can see that the PAIRS-STRICT and OKASAKI-STRICT performs a bit slower than their lazy counter-parts. This was expected but the difference between them is less than wat we anticipated. What is surprising is that HOOD-MELVILLE is slower than all of them (not regarding LIST which is THE slowest), but thus must have something to do with the constant of insertions and deletions that is for this example is 10.

7.3.5 Queue-Sort

This is the results of our real-world use-case implemented as a real-world (not computer-scientist) programmer would do:

Figure 7.5: Time of merging n queues with n elements into a single queue divided by n^3

When the number of queues and the elements in the queues increases, the running time of each of the queues goes towards the same point. This is because all elements must be removed from each queue, and there will for each queue sans HOOD-MELVILLE be a linear amount of expensive operations. But towards $n = 100$ we find that HOOD-MELVILLE and PAIRS-LAZY seems to be the fastest.

7.4 Summary on running times

In general all the running times are pretty close when we exclude the LIST implementation. This matches the theoretic bounds but we must admit we had hoped for bigger gaps.

One of the reason why we did not achieve might be that everything was compiled with an optimizing compiler, which closer matches what would be used in the industry (if the used functional programming). It would not remove function calls, but it would inline them making them faster.

Chapter 8

Functional lists with index lookup

This exercise was posted in the theoretical project but solving it seemed natural considering this project is all about functional data structures.

We are asked to describe a strict functional data structure supporting PUSH(x, L), POP(L) in O(1) and LOOKUP(d) in time O$(\log n)$ where n is the number of elements in the list.

By using a standard list composed by **cons** we will have a PUSH(x, L) and POP(L) in constant time, so how do we solve lookups in O$(\log n)$? A binary tree would give us the desired running times. The problem is how to combine the two. Okasaki solved this problem by having a collection of complete binary trees [3].

For each entry in our list we store a tuple consisting of a complete binary tree and the size the tree. Thus, the size of the trees is on the form $2^k - 1$. Whenever two adjacent elements are at the front of the list and another element is pushed, the trees are joined with the newly added element as the root which gives us that the elements in any of the binary trees is stored in preorder. Elements 1 to $\lfloor n/2 \rfloor$ will be in the left part of the tree and $\lceil n/2 \rceil$ to $n - 1$ will be in the right part.


Figure 8.1: A list consisting of three complete binary trees, two of size 3 and one of size 7. An added node would result in two complete binary trees of size 7 and an additional node would result in one tree of size 15.

Inserting new elements can at most join two trees. Since the left child of the resulting tree will be the first tree in the list and the right child will be the second tree, the tree can be constructed in constant time. Similarly, it can be destructed into two parts by just removing the root and **cons** the right child followed by the left child onto the list.

Looking up and element by index can be done by simply iterating the list until the containing tree has been found. If the index to search for is smaller than the size of the tree at the current position said tree would contain the element, otherwise subtract the size from the index and continue searching.

If we only have on complete binary tree looking up an element can clearly be done in $O(\log n)$ time so we have to show that the length of the list will never be longer than $\log n$.

An integer on the form $2^k - 1$ is called a skew-binary term, and if t is such a term then the next skew-binary term will be $2t + 1$. A decomposition of an integer n greater or equal to zero will be a multiset of skew-binary terms $\{t_1, t_2, \cdots, t_m\}$ where $n = t_1 + t_2 + \cdots + t_m$. Such a decomposition is said to be greedy if the largest term is as large as possible and the remaining part of the decomposition is greedy as well. In other words a collection of skew-binary terms is greedy if no subset of lesser terms sums to a larger term. We will describe such a greedy decomposition of n as $G(n)$.

Property 1. *Every integer $n \geq 0$ has a unique greedy decomposition.*

Proof. $G(0)$ will have the empty set which is unique. For $G(n)$ where $n > 0$, we know, that no subset of skew-binary terms may sum to a larger term thus the it must include the largest possible term t such that $t \leq n$. We add t to the set and continue searching for $G(n - t)$. Notice, that no subset of $G(n - t)$ can sum to t otherwise it would not be the largest term respecting $t \leq n$. $\square$

Property 2. *A skew-binary decomposition is greedy iif every term is unique except the two smallest: $t_1 \leq t_2 < t_3 < \cdots < t_m$.*

Proof.
$\Rightarrow$ We have a greedy decomposition and say we have two terms of size $2^k - 1$ and a third term of equal or lesser size. Then the three terms would be greater or equal to another term because the next term would be at least as large as $2(2^k - 1) + 1$. But then this would not be greedy and we have a contradiction.
$\Leftarrow$ Again, because the next skew-binary term for any term t is on the form $2t + 1$, no two terms can sum up to another term if we only repeat the smallest ones twice and all others are unique. The decomposition would therefore be greedy. $\square$

Theorem 1. $|G(n)| \leq \lceil \log(n + 1) \rceil$.

Proof. Clearly if $n = 0$ it holds so we turn our attention to $n > 0$. Assume, to obtain a contradiction, that $|G(n)|$ contains more than $k = \lceil \log(n + 1) \rceil$. Then, by 2 and by how a greedy decomposition is defined, it must be that $G(n)$ contains a term at least as large as $2^k - 1 \geq n$. Then, because $|G(n)|$ contains more than k terms there is at least one more term and therefore the sum of terms in $G(n)$ will strictly exceed n and we obtain a contradiction. $\square$

Except for the two first trees in our list every size of the complete binary trees in the list is unique and in increasing order by construction. Therefore, the length of the list can not exceed $\lceil \log(n+1) \rceil$ by 1. Therefore, lookup will traverse at most $O(\log n)$ elements in the list and perform a search in a single complete binary tree of height at most $\log n$.

Chapter 9

Conclusion

This is the last report in the course Advanved Data Structures at Aarhus University and for our last project we decided to work with functional data structures. We have shown four different implementations of queues.

When analyzing the queues time-complexities, we found that analyzing lazy evaluated structures required us to reformulate the potential function for amortized analysis.

We proved our implementations correct in Coq, which gave us the confidence to go ahead and actually perform experiments on them. Our experiments more or less showed that our theoretic bounds matched.

Finally, we wanted to show that it is possible to do $O\left(\log n\right)$ indexing with constant time push and pop, which was an exercise posted in the theoretical project.

Bibliography

[1] Thomas H. Cormen, Charles E. Leiserson, Roland L. Rivest, and Clifford Stein. *Introduction to algorithms*. The MIT Press, third edition, 2009.

[2] Robert T. Hood and Robert C. Melville. Real time queue operations in pure lisp. 1980.

[3] Chris Okasaki. *Purely Functional Data Structures*. Cambridge University Press, first edition, 1998.